

Traceable by design

Building a source-checked investigation with Claude Code

In five working days, one investigator used Claude Code to build a human-rights investigation in which every cited statement traces to a quoted source, delivered as a single offline file that refuses to build if any link breaks. This paper explains how it was done, what went wrong, and how a team can do it again.

1

operator making the
decisions

5

working days, Sep
24–29, 2026

165


prompts typed;
median 10 words

45

Claude Code
sessions, up to 6 at
once

473

subagent and
workflow launches

**1,813**

cited claims in the
finished report. One dot
each. Every claim carries
a verbatim excerpt from
a source that an agent
actually opened.

The short version

One investigator directed Claude Code to build an open-source investigation into alleged abuses in Sudan's Kordofan region: 1,813 cited claims, 277 incident profiles and 13 chapters, all in one file that works offline. Every cited statement traces to a claim, and every claim to a quoted source.

WHY IT MATTERS

Anyone can check it. A reader can follow any cited statement to its claim, and any claim to the words on the source page, without internet access.

Checked by design, not by promise. Agents that wrote never checked their own work. Scripts and separate, skeptical agents did.

Reusable. The method is now a kit with no Sudan-specific data, ready for the next investigation.

WHAT IT TOOK

One person deciding. 165 prompts, most of them a single line: a choice, a rule or a screenshot.

Machine effort. 45 sessions, 473 agent launches, about 6,800 page fetches, 25.0 million output tokens.

Five working days from the first prompt to the finished file and a reusable kit.

WHAT IT DOES NOT PROVE

That the sources are right. The checks prove that claims can be traced to sources, not that the sources are true. An expert reader is still needed.

That it beat a team. Nobody ran the same project without Claude, so there is no speed-up figure.

What it cost. Token counts are recorded; a dollar figure was not calculated (\$8).

The project worked because the operator spent effort on constraints, not instructions: what counts as evidence, who may be named, what a build may never do.

THE CENTRAL FINDING

Six lessons for anyone trying this

01 Start with an interview.

Have Claude question you before it researches. Four answers (time window, actors, evidence bar, naming rule) shaped everything after.

03 Separate the writer from the checker.

Research agents drafted; different agents, told to assume errors, verified. They changed a large share of claims before anything was published.

05 Ask for a menu before you ask for work.

"What gaps exist? Don't implement anything."
Then pick by number. One line per decision keeps a person in charge.

02 Make the evidence model the product.

Prose cites claim IDs, claims carry quoted excerpts, and the build fails on any unknown ID.

04 Agents propose, code applies.

Agents never edited shared files. A script backed up, validated and applied their proposals.

06 Write rules where Claude will re-read them.

A global `CLAUDE.md` and about 30 memory notes carried standards across 45 sessions.

What a team needs to try this

-
- ☐ **One operator with authority** over scope, evidence bar and naming, free for short decisions over about a week.
 - ☐ **A capable machine with a set agent cap**, plus quotas for search, social-media and satellite data.
 - ☐ **Protected-data rules**, written before research starts: who may never be named, and where files may live.
 - ☐ **A release gate that isn't the operator**: a privacy scan, a fail-closed checker and an outside critique.
 - ☐ **A subject-matter expert** to read the result. The checks prove traceability, not truth.
-

The sixteen incidents that shaped the method are in §6; the full playbook is in §7.

Before you start

The rules this paper follows, and the terms it uses.

Conventions

Quotations are verbatim. A build script pulls prompts and Claude's replies from the session transcripts and fails if any quotation can't be found. The operator's typos are left in. Cuts are marked `[...]`, and Claude's emoticon sign-offs are removed.

This paper is about method, not findings. Personal names are replaced with `[name]`, and so are details that would reveal a name. Places, units, incident IDs (INC-####) and file names are kept. No death tolls or allegations against named people appear here.

Labels. Times are the operator's local time (UTC-7). Codes like "S07" are sessions, listed in Appendix A. "Claude" means Claude Code, which ran Claude Opus 5.5 on nearly every turn and Sonnet 5.5 on a few hundred.

Terms used in this paper

Operator

The person directing the work. Sets scope and rules, makes decisions, reviews results.

Workflow

A script that launches many subagents in phases, for example research first, then verify.

Fail-closed

A build that produces nothing if any check fails, instead of warning and carrying on.

Session

One Claude Code conversation working in the project folder. Several can run at once.

Claim

One sourced statement (who, what, where, when) with a grade and a verbatim excerpt of up to 25 words, which a script checks against the source page. Chapters cite claims by ID.

Fan-out

Launching many agents in parallel, each on one slice of a task.

Subagent

A helper that a session launches for one task, such as researching one topic. It sees only its written brief.

Grade A / B / C

Confidence in a claim, based on how many independent, credible sources support it (Appendix C).

Memory note, `CLAUDE.md`

Short files that Claude re-reads in every session. They are how rules outlive a conversation.

The project

What the investigation asked, the rules set in the first hour, and what the finished product contains.

1.1 The ask

On September 24, 2026, the operator asked Claude Code for a fact-only investigation into alleged abuses by the Rapid Support Forces (RSF) and allied groups in Sudan's Kordofan region. The first prompt, below in full, named a role, a subject, the entities to find, a tone, a citation format and a guard against invented facts, and it asked Claude to ask questions. It said nothing about technology: no framework, no file layout, no tools.

THE OPENING PROMPT

Sep 24 · 00:12

S01

> I want you to create a report from the perspective of an human rights investigator examining the RSF's fighting in the Kordofan region. I want you to identify units, commanders, specific people, specific incidents of human rights abuses, mass killings, torture, rape. This initial report should be focused purely on the facts. I want you to create a html report, citing all facts as footnotes. I want you to review the work to check for factual correctness and to eradicate all hallucinations. Ask me questions to design your investigation and research. Orchestrate as many agents as you need to do the work thoroughly, quickly and effectively

1.2 Rules set in the first hour, and never relaxed

DECISION	CHOICE	WHY IT MATTERED LATER
Time window	2025–2026 only	Bounded the research rounds and made “thin months” measurable.
Actors	RSF and allies; the army only as context	Defined “in scope” for every incident and every later audit.
Evidence bar	Every claim graded A, B or C; partisan and social-media sources usable only as “party X says”	Grades drove the corroboration passes and limited what the cover could claim.
Naming	Name alleged perpetrators; never name survivors of sexual violence or torture; name the dead only where a cited source does	A privacy scan and a rule excluding weak incidents both grew from this line.
Vocabulary	Descriptive categories only; no legal labels such as “war crime”	Kept the product factual and defensible.

DECISION	CHOICE	WHY IT MATTERED LATER
Accuracy stance	"Be more conservative... do not reach"; show discrepancies, never hide or overstate them	The tie-breaker for hundreds of later judgment calls (§4.3).

1.3 What was built

The deliverable, called the *edition*, is one self-contained 19 MB HTML file. It has three layers. Each rests on the one below, and a checker must pass before the file exists.




FIGURE 1 Anatomy of the edition. Counts are from the final checker run; the source and publisher counts are from a Sep 26 audit.

WHY A SINGLE FILE

One offline file means nothing to host, nothing to break, and one thing to hand a colleague or auditor. It also made a strict build possible: the builder inlines everything, and the checker validates the whole file before release. A companion tool can encrypt it behind a password.

Five days, five phases

Activity peaked in the first two days, when scope, evidence and design were decided. The later days refined, audited and distilled.




FIGURE 2 Activity per day, from the transcripts (local time). Agent launches are those made by the main sessions; the agents' own tool calls are not counted.

PHASE 1

Sep 24
00:12–08:30
S01

Scope and a first evidence base

Scoping questions, a nine-track research workflow, a sweep of Arabic-language and partisan sources, and a “completeness critic” agent that listed gaps. A `/goal` kept the session going until sources were “reasonably exhausted”. Result: about 1,100 verified claims and a footnoted HTML report.

PHASE 2

Sep 24
11:00–17:00
S02–S08

Legibility, people, integrity

The report had grown too long to navigate, so it was rebuilt as an app. In parallel, one session added photos and Arabic name forms, one built a biography for every named person, and one audited gaps and regraded claims. The day ended with a pause and a written hand-off.

PHASE 3

Sep 24
17:30–21:00
S09–S12

Analysis and the incident database

The analysis page was redesigned for clarity. The incident database was planned first, approved in one line of decisions, then run in stages. Unfinished satellite work was parked, not shipped.

PHASE 4

Sep 24 22:50 –
Sep 25 23:45
S13–S36

The edition: redesign and layers

The busiest stretch: 21 sessions in one day. A ground-up redesign produced the single-file edition. Separate sessions then added satellite analysis, order of battle, unit profiles, responsibility views, background, social-media evidence, an encryption tool and a counts audit.

PHASE 5

Sep 27–29
S37–S45

Improve, close gaps, distill

Review-and-pick rounds improved the edition. Telegram channels were crawled and analyzed, research gaps were closed, and the process was distilled into a reusable kit. This paper was written in the last session.

How it worked

Eight mechanisms carried the project. None of them is specific to Sudan or to human-rights work.

3.1 Interview before building

Told to “ask me questions to design your investigation”, Claude used its `AskUserQuestion` tool to offer four multiple-choice questions, each with a recommended option and its consequences. The answers became the scope note, and later a dated `DECISIONS.md` in the kit. §5.1 shows what happened when the operator closed the dialog without answering.

QUESTION	OPTIONS OFFERED (RECOMMENDED FIRST)	OPERATOR'S CHOICE
Time scope	Whole war · 2025–2026 only · a specific window	2025–2026 only
Actors and geography	RSF + allies, all three states · RSF only · RSF + army	RSF + allies, all three states
Evidence bar	Tiered and graded · corroborated only · include OSINT and social media	Tiered, with OSINT and social media
Naming	Name perpetrators, not victims · commanders only · everyone in sources	Name perpetrators, never survivors

3.2 The evidence pipeline

The process that emerged, later written down as a 17-stage kit, alternates three kinds of work: *agents* produce, *scripts* check mechanically, and the *operator* decides at gates. At a gate, nothing continues until a check passes or the operator decides.

THE THREE RULES INSIDE THE PIPELINE THAT MATTERED MOST

- **Every claim carries a verbatim excerpt of at most 25 words** from a page an agent actually opened, and a script re-checks the excerpt against the live page.
- **The verifier is a different agent, told to assume every claim may be invented.** It may keep, correct or drop a claim. It can't rescue a source that failed the mechanical check.
- **High-stakes claims get a red-team agent** whose only job is to disprove them. In practice this mostly softened how strongly conduct was attributed to named individuals.




FIGURE 3 The evidence pipeline. Stages 0–11 establish the facts; 12–16 build and maintain what people read.

3.3 Agents propose, code applies

Parallel agents never edited shared files. Each wrote one JSON file to its own folder and replied “done”. A script then backed up the target, validated each proposal, flagged anything suspicious (a verdict flip, a large deletion, a grade change) and applied the rest. Scripts defaulted to a dry run and needed `--apply` to write. The project built up 23 backup folders this way, each a known-good state to return to.

3.4 Builds that fail closed

Chapters cite claims with markers such as `{{c:ID}}`. The builder turns markers into footnotes and stops if an ID is unknown or rejected. A prose checker flags any number in the text that doesn't appear in the cited claims. The edition checker then confirms that every incident, person and claim link resolves and that the inventory is complete. The final run printed 13/13 sections, 1,813/1,813 claims, 277/277 incidents and “PASS”. If a check can't pass, there is no file.

3.5 Many sessions, one codebase

From the first afternoon, the operator ran several Claude Code sessions at once, each with its own focus. Sessions could message each other and were notified when a session they were waiting on went idle. The operator's instruction was one line: *“there are two other claude sessions working on this project right now, please coordinate with them”*. The sessions then agreed file ownership and a build lock among themselves (§5.2).




FIGURE 4 Sessions coordinate through files and messages (dashed lines: messages between sessions). Ownership was agreed in chat, recorded in memory notes, and revised when a session finished.

3.6 Goals that keep a session working

For open-ended work the operator used `/goal`, which stops Claude from ending its turn until a stated condition holds. Four goals were set (Appendix B), each a finish line in the operator's words: “reasonably exhausted the available open source information”, “production ready and can be published to a wide audience”, “deep QA on the tool to prepare it for use by investigators”. They let the operator walk away.

3.7 Memory as durable state

A session forgets; files don't. Two mechanisms carried standards from one session to the next:


- A **global** `CLAUDE.md` held rules for every project: the agent cap, warnings about workflow pitfalls, and style preferences. Claude edited it when the operator asked, for example to raise the cap.
- **About 30 project memory notes**, each a short file with a type, a one-line description, and a *why* and *how to apply*. A session starting cold reads them and applies the rule correctly.

EXAMPLE: HOW A RULE BECAME A MEMORY NOTE

Mid-task, the operator said *"use your best judgment to address those notes. The priority is accuracy ... do not reach"*, and later *"you should highlight discrepancies, don't hide things but don't overstate"*. Claude saved both as one note. Conflicting sources are shown neutrally and attributed. Weak text is narrowed, not extended. Grades are lowered only when in doubt. Judgment calls are resolved and reported rather than sent back as questions.

3.8 The human's job: decisions, not dictation

Most of the operator's prompts were decisions on options that Claude had prepared. This loop repeated in every improvement round:



The prompts

33 of the operator's prompts, verbatim, grouped by purpose, with a note on why each group worked. The complete log is in Appendix B.

THE OPERATOR'S PROMPTING STYLE

- **Short.** The median prompt was 10 words; 78% were under 25.
- **Outcome first.** Prompts described what the reader should experience or what must be true, not how to code it.
- **Trusting, with brakes.** "Use your best judgment" appears often, always alongside a standing rule on accuracy, scope or naming.
- **Visual.** Ten prompts came with a screenshot. A picture plus one sentence was the fastest bug report.

4.1 Frame the work

Give Claude a role, a standard of truth and a finish line. Let it ask for what it needs.

RECOVER A DISMISSED DIALOG

Sep 24 • 00:13

S01

> ask me the questions again please

WHY IT WORKED The operator had closed the first dialog without answering, and Claude correctly paused rather than guessing. One line reopened the interview. Never let Claude fill in scope answers itself.

EXHAUSTIVENESS AS A GOAL

Sep 24 • 00:20

S01

/goal condition

> make sure to look at all sources including arabic language sources and sources from the SAF and RSF sides. keep working until you feel like you've reasonably exhausted the available open source information

WHY IT WORKED It defines "done" by coverage (all languages, both sides) and gives permission to keep going. The Arabic and partisan-source requirement produced a whole second research round.

DESIGN IN THE SAME BREATH AS CONTENT

Sep 24 · 07:54

S01

- › I want you to add an always present ToC to the .html report to help the user navigate the report. I also want you to add an elegant, professional and beautiful design to the .html file. Use charts and tables where relevant to give users a high level view of the information.

WHY IT WORKED Quality words (“elegant, professional and beautiful”) and a use (“a high level view”) were enough for a first pass. The design was rebuilt within hours, once the content outgrew it (§4.4).

AMBITION WITH A LEGIBILITY LIMIT

Sep 25 · 08:48

S19

- › I'd like for you to add a dedicated satellite imagery analysis section in the Explore the Evidence section and build it out. The goal should be to use sentinel 1, 2 and wayback imagery to illuminate the incidents. Use your best judgment to determine which incidents can be explored using the available data. I want you to push the boundaries a bit here, but keep it legible, the analysis understandable. use GEE and other tools as needed.

WHY IT WORKED “Push the boundaries a bit... but keep it legible” invites exploration while protecting the reader, and Claude decides which incidents the imagery can support.

4.2 Menus first, then picks

The most reusable pattern: ask for options without work, then choose by number. Use it whenever the next step is unclear.

ASK FOR A GAP LIST

Sep 24 · 11:38

S06

- › What gaps exist in the report and how can we improve it?

MENU WITH AN EXPLICIT STOP

Sep 25 · 17:34

S28

- › how can we either further visualize and analyze the existing data or expand our research? What investigative gaps exist that we might be able to address? do not implement anything I am just asking

WHY IT WORKED “Do not implement anything I am just asking” prevents a well-meant sprint. Claude replied with a table of measured gaps (§5.7), ranked by value for effort.

PICK BY NUMBER, DELEGATE THE REST

Sep 25 · 19:31

S28

- › use your best judgment to address the open decisions. I want you to implement unit attribution and the easy work first. do 1, 2, 5, 6, and 7

WHY IT WORKED The numbers refer to the menu just shown. Claude noticed they could match two different lists and checked which was meant. Keep menu numbers unique within a session.

THE SAME LOOP, DAYS LATER

Sep 28 · 18:12

S42

> how can we improve the edition version?

Sep 28 · 18:15

S42

> let's do 1 - 4

Sep 28 · 18:36

S42

> move those 5 out of scope

WHY IT WORKED “Move those 5 out of scope” records a decision in the project's own scope language, and it closes the loop on items you won't do.

PLAN FIRST; DO NOT IMPLEMENT

Sep 24 · 19:14

S11

> I'd like for you to plan and design a process to enrich the current incidents we have with further research to build out incident profiles and a database of incidents in Kordofan that could be a human rights abuse or atrocity. Do not implement the work until I give the go ahead

ONE LINE ANSWERS SEVEN DECISIONS

Sep 24 · 19:19

S11

> Descriptive only, RSF+allies, parent campaigns, new tab, feeding back, name killed victims. unleash the search budget, search as much as you need.

WHY IT WORKED Claude returned a plan with open decisions. The operator answered all of them in one comma-separated line and lifted the budget (“search as much as you need”). Claude recorded the decisions, then ran the cheap steps before spending on agents.

4.3 Set standards and enforce them

These prompts state a rule, not a task. Most became memory notes.

THE ACCURACY RULE

Sep 25 · 10:16

S22

> I want you to use your best judgment to address those notes. The priority is accuracy - be more conservative with the information if things are hard to determine from the open source - do not reach

WHY IT WORKED It tells Claude how to resolve ambiguity (narrow, don't reach) and hands over the judgment. It is the project's philosophy in one line, and it is now in the kit's operating rules.

SPELLING IS DATA

Sep 24 · 14:56

S07

+ 1 screenshot

- › the report should use names that reflect the latin spelling found in reports. like in this example, it should not transliterate in a way that is not represented in the cited sources

WHY IT WORKED One screenshot of one wrong transliteration became a rule (use only spellings the cited sources print), a mapping file, and scripts that applied it across the report.

KEEP IT IN THE DATABASE, OUT OF THE REPORT

Sep 24 · 21:02

S11

- › leave out weak new incidents, keep it in the database. note that there are weak incidents

WHY IT WORKED It separates *storing* from *publishing*. Weak material stays searchable and flagged, but doesn't shape the report.

SPOT-CHECK BY NARRATION

Sep 25 · 09:22

S20

- › first walk me through INC-0154. then address INC-0129 and the Misfiled claims.

...THEN ACCEPT THE RECOMMENDATION

Sep 25 · 09:25

S20

- › i'll take your recommendations here, use your best judgment and implement fixes for all three

WHY IT WORKED "Walk me through" makes Claude lay out the evidence, and that is how the operator caught a scope error (§5.4). Accepting the fix took one line.

4.4 Redesign and feedback

Design prompts described the reader's experience and named a standard. Screenshots supplied the specifics.

PERMISSION FOR RADICAL CHANGE

Sep 24 · 11:07

S03

- › this report is so long now it's hard to read. the ToC breaks when I jump to certain sections because there is too much content. Please redesign this report for legibility and clarity. If sections need to be a database or a radical redesign, do so

WHY IT WORKED "If sections need to be a database or a radical redesign, do so" lifted the constraint that would have kept Claude patching a document that had outgrown its format. It rebuilt the report as an app with a filterable incident explorer.

THE GROUND-UP REDESIGN

Sep 24 · 23:30

S15

- › I want you to redesign file://~/sudan-research/report/kordofan_rsf_package.html from the ground up. it was made in an interative process and is not coherent or unified. I want you to recreate it, keeping the facts and information, in a way that helps the reader understand the situation, who is responsible, and be able to easily explore the underlying data. It should balance narrative with data exploration. It should follow best practices like those outlined by Tufte, but do not feel constrained by it. Strive for beauty, elegance, legibility, understandability. Help the user understand the information - show linkages, patterns and highlight connections. Make this a delight for the user to read and experience but keep it informative - something the user can come back to to use as a tool. Use your full creativity

WHY IT WORKED It gives a diagnosis ("not coherent or unified"), a constraint ("keeping the facts"), three reader goals, a named standard (Tufte, "but do not feel constrained by it") and room to be creative. Claude designed, had an independent agent critique screenshots, then refined once. The critique found integrity problems (\$5.8).

SAY WHAT'S WRONG, KEEP WHAT'S GOOD

Sep 24 · 18:13

S10

- › I'm not happy with the how the analysis tab is designed. I find it hard to understand. Please rework it so that the analysis is clear and able to be immediately understood. Keep the nuance.

WHY IT WORKED "Keep the nuance" stops simplification from turning into overstatement. The result was a fixed pattern: short answer, graphic, caveats, detail.

BORROW A PATTERN FROM ELSEWHERE

Sep 25 · 22:54

S33

- › I want you to find a better way to represent the information contained in the Responsibility section. style each named commander and allow for grouping and sorting in a similar way that the "The Three Fronts" section organizes its information. Take inspiration and localize it to make sense for the Responsibility section

WHY IT WORKED Pointing at a design the reader already knows ("The Three Fronts") is faster and more precise than describing one from scratch, and "localize it" rules out a literal copy.

SCOPE A PASS BY WHAT NOT TO TOUCH

Sep 25 · 23:29

S35

/goal condition

+ 1 screenshot

- › I want you to examine all charts labels and relevant sections to ensure that the labels are clear, simply written and legible. For example, the labels and titles in this screenshot do not read well and feels overly complicated to read. You are only adjusting the text, not the design or graphics themselves

WHY IT WORKED "You are only adjusting the text, not the design or graphics" confines a long-running session to a safe layer while other sessions edit the visuals.

A SCREENSHOT AND ONE SENTENCE

Sep 25 · 14:09

S27

+ 1 screenshot

- › when when i select in the chart it doesn't seem to actually all appear in the map

WHY IT WORKED Claude reproduces the problem from the picture. Screenshot prompts fixed a dead profile link, a chart-to-map mismatch and confusing pattern counts, among others.

4.5 Parallel sessions and pacing

The operator managed capacity and attention, not code.

HAND COORDINATION TO CLAUDE

Sep 24 · 11:40

S06

- › there are two other claude sessions working on this project right now, please coordinate with them

RAISE THE CAP, WRITE IT DOWN

Sep 24 · 14:05

S06

- › we can use more than 2 agents at a time. let's set the limit to 8 total

WHY IT WORKED The operator stated the situation and the limit, not a protocol. The sessions worked out ownership and a build lock themselves (§5.2), and the new cap went into the global `CLAUDE.md`, so every session followed it.

SEQUENCE WORK ACROSS SESSIONS

Sep 25 · 22:05

S30

- › I want you to redesign the cards in the north, south and west kordofan sections to be organized in the existing timeline format or with a toggle option to be organized into groups based on categories of types of attacks. there are two other sessions working in those codebase. one is working on the cards right now, coordinate so that you don't conflict. wait until it's done with updating the cards to implement any changes. while you wait, think of how you can make this design elegant beautiful and useful

WHY IT WORKED It names the conflict ("one is working on the cards right now"), says to wait, and puts the wait to use: "think of how you can make this design elegant".

BAN IDLE WAITING

Sep 25 · 10:06

S22

- › I want you to be proactive and work on other things if you are able to

WHY IT WORKED Agent rounds can take many minutes. This standing instruction, saved as a memory note, turned waiting into useful work.

PAUSE AND HAND OFF

Sep 24 · 16:20

S08

- › actually let's pause, I'm going to relocate

Sep 24 · 16:21 S08

› create a handoff for the next session

WHY IT WORKED Claude saw the real risk (its scratch folder might be cleared), copied the working files into the project, and wrote a resume script so the next session could finish with one reviewed command (§5.6).

4.6 Sources, scope and the kit

UNBLOCK A KEY SOURCE Sep 24 · 15:22 S07

› can you find a way around sudan tribune's blocked automated fetching? i think they're a key source

ESCALATE TO THE OPERATOR'S OWN BROWSER Sep 24 · 15:26 S07

› use my chrome

WHY IT WORKED Claude tried every automated route first (§5.5), then proposed the least invasive option: read the public articles in the operator's own browser, read-only, as any reader would. Three words approved it. A step like this should be a recorded decision.

ASK FOR AN ESTIMATE BEFORE COMMITTING Sep 28 · 11:38 S41

› how much work would it be to add a section looking at SAF activities? I want to keep it separate from the RSF analysis but still in this edition version

WHY IT WORKED It asks for size before commitment and states the one constraint that matters: keep it separate from the RSF analysis.

TURN THE PROJECT INTO A KIT Sep 29 · 09:49 S44

› I'd like for you to distill the research and production process of the edition report into a reusable process. The research should determine the specific structure of the report but this is the general skeleton that I like. Make this into a private repo called auto-researcher

WHY IT WORKED It says what stays fixed ("the general skeleton that I like") and leaves the rest to the next project's research. Claude separated the Sudan-specific parts from the reusable ones and produced the 17-stage kit.

Inside the conversations

Eleven moments where the method was made, tested or broken. Claude's words are verbatim; names are redacted.

5.1 The dialog that was closed

WHAT TO NOTICE

Claude didn't go ahead on its own defaults. Scoping was a hard gate, and getting past it cost the operator one line.

● CLAUDE

S01 · Sep 24 · 00:13

You closed the scoping questions without answering, so I haven't started any research. Nothing has been searched or written, and the project folder is still empty.

When you're ready, send your choices in chat or tell me to go with my suggested defaults.

[...]

Sep 24 · 00:13

S01

> ask me the questions again please

5.2 Sessions negotiating who owns what

WHAT TO NOTICE

Told only to “coordinate”, Claude proposed ownership, a build lock and an order of work, and told the other sessions what it would leave alone. None of this was scripted by the operator.

● CLAUDE

S06 · Sep 24 · 11:41

I messaged both sessions that are working on the project with me: **Images for named individuals** and **Individual evidence base profiles**. I haven't edited anything yet and am waiting for their replies.

What I told them:

- **What I'd change:**
 - The citation fix in `report/sections/sexual-violence-torture.html`.
 - Standardising publisher names and adding an `originating_source` field in `data/verified*`. I'd back those files up first and apply the change with a script, and some grades may change as a result.
 - Later, an incident map in the chart and page code.

- **What I'll leave alone:** the photo, bio and profile folders, the code that builds people cards and profiles, and the section on named individuals. Those look like theirs.
- **What I asked them:** which files they're writing to, and whether they read the claim files. The profiles session probably shows claim grades, so a re-grade could change what it shows.
- **What I proposed:** only one session runs the report build at a time, and it tells the others before it does. The build writes to the same report files, so two builds at once would conflict.

[...]

5.3 Claude retracts its own finding

WHAT TO NOTICE

A detector flagged a missing citation. Claude checked, and reported that its own flag was wrong. Detectors produce false positives, so the kit's rule is to verify every hit before fixing anything.

● CLAUDE

S06 · Sep 24 · 11:41

I was wrong about the citation error, and I'm taking it back. The paragraph introducing the list already cites `gap_01_un_documents-017`, which contains every one of those figures, and it says they are national counts for 2025. `check_prose.py` linked the list items to the marker that comes *after* them, so the flag was a false positive. The section doesn't need any change.

[...]

5.4 "Walk me through INC-0154"

WHAT TO NOTICE

A request for a walk-through exposed a classification bug: a parser had read "RSF-held" as "by the RSF". Claude first laid out the incident's evidence (omitted here). Below: the cause, a check of similar cases, and a one-line recommendation.

● CLAUDE

S20 · Sep 25 · 09:24

[...]

How it ended up in scope. When the database was first built, a program sorted each card by the "Perpetrator" line printed on it. This card's line says the areas were "RSF-held", and the program apparently read that as "carried out by the RSF", which put the incident in scope. I checked the other in-scope cards whose perpetrator lines mention an unnamed or unknown unit; INC-0154 is the only one with no RSF attribution at all.

Recommendation: mark INC-0154's attribution as "not stated" in the incident registry. That's the same treatment 20 other incidents already have. It would:

- take it out of scope, so the cover would show 218 in-scope incidents instead of 219;
- leave the card and its wording unchanged;
- [a third effect, on a headline count, omitted].

The Um Nabaq strike (RSF per local sources) could become its own small in-scope incident, but it has no deaths. Should I make the INC-0154 change?

[...]

5.5 A blocked source, and the limits of “found it”

WHAT TO NOTICE

Claude tried the site's API, feeds, archives and the in-app browser, then stopped. A search-engine summary is not confirmation. It proposed the least invasive next step, read-only, and asked first.

● CLAUDE

S07 · Sep 24 · 15:25

The in-app browser is almost certainly being flagged as automated, so Cloudflare keeps re-asking no matter how many times you click. That can't be fixed from our side.

A search-engine summary does show both names on the two blocked pages: [article A] ("[name]") and [article B] ("[name]"). But that's the search engine's summary, not the page itself, so I don't count it as confirmation. No other outlet republishes these two stories. [Darfur24](#) covers the same event as [article B] but doesn't name [name].

Two options that should work:

1. **I use your own Chrome (Claude in Chrome)** (recommended). Your normal Chrome should pass Cloudflare like any ordinary visit. I'd open each Sudan Tribune article in a new tab and read the text; I'd only read, never click or type anything. After these two, I'd work through the other Sudan Tribune sources that were only confirmed indirectly.
2. **You open the pages yourself.** Open the two links in your own browser and paste me the paragraph that names each person. That's fine for two articles, but it doesn't scale to the rest.

Should I go ahead with option 1?

5.6 “Create a handoff for the next session”

WHAT TO NOTICE

Before writing the hand-off, Claude found the real risk (temporary scratch files) and built a resume script, so the next session needs one reviewed command.

● CLAUDE

S08 · Sep 24 · 16:22

The main risk for the next session is that everything in this session's scratch folder under `/private/tmp` could be cleared: the review decisions, the redo agent results, the command-responsibility reviews and the companion proposals. I'll copy it into the project first, then write the handoff against those permanent paths

[...]

Everything is copied into `analysis/phase3_work/` (5.4 MB). Next, a resume script, so the next session can finish with one reviewed command rather than reconstructing the steps

[...]

5.7 What a good menu looks like

WHAT TO NOTICE

Each gap is sized from the database and says why it matters. That is what lets the operator decide in one line.

[...]

Gaps the data itself shows

These come straight from `data/incidents/incidents.json`, counting in-scope incidents only:

GAP	SIZE	WHY IT MATTERS
No perpetrator unit found	128 of 218 incidents marked <code>perpetrator_unit: none_found</code> ; only 65 have command links	This is the biggest weakness for accountability. It's usually "RSF", with no unit or commander.
Sexual violence and child recruitment	6 incidents each, against 56 drone strikes	Almost certainly under-reporting rather than absence. These crimes seldom show up in press-led collection.
Few sources behind most of it	Darfur24 is in 93 incidents and Sudan Tribune in 91. 36 incidents have ≤1 non-partisan origin, and 60 are grade C at best	If one monitor has a blind spot or bias, the whole picture inherits it.
Place is vague	103 incidents are "near" and 19 have no location	This limits satellite checks and any spatial analysis.
No victim names	72 incidents are <code>none_found</code> ; 56 report deaths but name nobody	This is the memorialisation and accountability gap.
Counterclaims barely searched	148 are <code>none_found</code>	Worth knowing if the report will face scrutiny from the RSF side.
Thin months	Sep 2025 has 2 incidents; Jan, Mar and Apr 2025 have 3–5 each	These could be real lulls, the rainy season, or collection gaps. We haven't tested which.
Out-of-scope date	INC-0002 starts in August 2023	Scope is 2025–26. It's probably fine as a continuing incident, but worth a deliberate decision.

[...]

5.8 The independent design critique

WHAT TO NOTICE

The critique agent was told to find problems, not praise. Its top finding was about *integrity*, not looks: a visual encoding that overstated certainty. Afterwards every encoding was checked against one question: does this overstate certainty?

● CRITIQUE AGENT

edition/CRITIQUE.md

Single biggest weakness: integrity leaks in the encodings. Toll circles are sized by the *highest* claim. 138 incidents that do have killed figures are drawn as "no toll". Composite grades collapse to their best letter. Placeholder start dates are shown as exact days. In an accountability product each of these overstates certainty, and each is fixable.

5.9 Agent capacity is a shared resource

WHAT TO NOTICE

Raising the cap from two to eight led straight to arbitration: who gets the next free slot, in what order, with a promise to message each session.

● CLAUDE

S06 · Sep 24 · 14:06

[...]

Running now: 7 corroboration agents (batches 1–7, 122 claims) plus the profiles session's 1 agent, for 8 in total. Batch 8 (9 claims) is still queued.

Both other sessions want more slots, and the machine is full. I've queued them in the order they asked, so nobody goes over 8:

1. The first slot my batches free goes to the profiles session, which has two independent biography batches to run.
2. The next two go to the images session, which you've asked to find photos for 53 people.
3. My batch 8 runs after that.

I'll message each session when its slot frees.

When results come in: I'll run each finished batch through the checking script in dry-run mode. You'll see every proposed grade change, with its reason, before anything is applied. Any source contradicting a claim is flagged for your review, not applied automatically.

5.10 An instruction the platform couldn't honor

WHAT TO NOTICE

Claude said plainly what it couldn't do, then changed the design so the answer would still exist afterwards: agents record their advisor questions in their own output files.

● CLAUDE

S43 · Sep 28 · 23:39

No, the advisor can't relay those questions to me. Each advisor call sees only the transcript of whoever calls it. When a subagent asks the advisor something, I never see the question or the answer, and there's no shared log.

I've done the next best thing. I messaged all four running agents. Before they finish, each will record whether an advisor tool was available to it, every question it asked, and the advice it got. The three geocoders will put this in an `advisor_notes` field in their proposal files. The sweep agent will put it in an "Advisor" section of `NOTES.md`. I'll read those back and summarize them when the agents finish.

I also don't know yet whether the subagents have the advisor tool at all. If they don't, your "let them ask a fable agent" instruction wasn't met for this batch. In that case I can re-run the ambiguous geocoding cases through an agent with `model: "fable"`, if you want that.

● CLAUDE

S43 · Sep 28 · 23:42

A safeguard error means an automated content filter on the API refused to let the model continue. It isn't a bug in our code or a problem with the task. The message in the failure said the Sonnet 5.5 safeguards "flagged this message" under Anthropic's acceptable use policy. It also said this "sometimes happens with safe, normal conversations", meaning it can be a false positive.

I don't know what triggered it here. The details line pointed at "reasoning_extraction", which suggests the filter reacted to something in how the agent was reasoning or being prompted, rather than to the geocoding task itself. Two identical Sonnet agents failed and the Opus reruns are running, so it looks specific to Sonnet 5.5 on these prompts. That's a guess, not something I've confirmed.

The practical effect is that the failed agents wrote nothing, so no data was lost, and I reran their jobs on Opus.

[...]

5.11 A link that broke only for the reader

WHAT TO NOTICE

Claude's tests passed on a local web server, but the operator's screenshot showed every internal link dead when the file was opened from disk. Claude explained why its tests missed it, fixed it, and said plainly which case it still hadn't tested.

● CLAUDE

S06 · Sep 24 · 14:55

[...]

The cause. Opened from your disk, the file has a `file://` address. The package loaded each tool from a temporary `blob:null/...` address, which is what your browser's status bar showed. Chrome won't follow in-

page `#/...` links inside such a document from a local file. So every link inside the tools was dead: sidebar names, links mid-page and footnotes. My first round of testing used the local server, which is why I missed it.

The fix. There are no temporary addresses any more. Each tool is written straight into its frame, and the package handles every link inside the tools itself. Only the package's own address and history change, which is the same mechanism your working tab bar already uses. The one piece of code that changed the report's address directly, pressing Enter in the search box, now goes through the same handling.

What I checked on the local server:

- **Sidebar and in-page links:** a sidebar name opened [name]'s profile, and a footnote link inside the profile opened the right reference.
- **Between tools:** a "Report claim" chip opened the evidence register filtered to that one claim, and a link inside the report then opened "At a glance".
- **Back and Forward:** Back four times retraced every step to the starting page, and Forward repeated them. Each click added exactly one history step. No tool ever went blank or loaded the whole package inside itself; an earlier attempt at this fix failed exactly that way.
- **Also fixed:** external source links now open in a new browser tab instead of replacing a tool inside the package.
- **No console errors.**

If anything still doesn't work, tell me which link and which tab.

If you'd like me to test the file-on-disk case myself, connect Claude in Chrome: install [the extension](#) and sign in from its side panel.

What went wrong, and the rule each produced

Each rule below was paid for by an incident in this project. The incident is the reason to believe the rule.

AREA	WHAT WENT WRONG	RULE ADOPTED
MACHINE Fan-out	Four sets of workflow agents launched at once froze the operator's computer.	Cap concurrent agents across the whole machine , not per session. Raise the cap in steps and lower it after any freeze. Here it went 2, 8, briefly 4, back to 8, then 12.
MACHINE Hung agent	A research agent using a browser tool hung for hours; a download with no timeout held a slot.	No browser tools for research agents. Every download gets a timeout. Check for stragglers.
BRIEFING Refusal	A workflow agent refused its task when its only context was "ask me the questions again".	Every agent brief opens with <i>who asked, why, and that the task is authorized</i> , in the user's own words (Appendix C).
OUTPUT Lost results	Agents asked to return large structured results often never returned them.	For bulk work, each agent writes one plain JSON file and replies "done". Use structured returns only for small results.
INTEGRITY Backups	Backup files left in the data folder were loaded after the live files and silently reversed a regrade of 70 claims.	Never store backups where the loader reads. The loader now skips backup files and warns.
INTEGRITY Duplicates	The same article under a differently encoded URL was counted as two sources.	Normalize URLs before de-duplicating.
INTEGRITY Independence	Corroboration agents called sources "independent" too generously ("according to Sudanese media" confirms the event, not the figures).	Review every proposed upgrade by hand. Grade A needs two independent, non-partisan origins. A weaker corroboration never lowers a grade.
INTEGRITY Classification	An incident entered scope because "RSF-held" was parsed as "by the RSF" (INC-0154).	Ask for walk-throughs of edge cases. Look for the same pattern elsewhere before fixing one instance.
INTEGRITY Encodings	Circle size showed the highest reported figure; 138 incidents with figures were drawn as "no toll"; placeholder dates looked exact.	Get an independent critique before release. Test every visual encoding for overstated certainty. Never add up per-claim figures.

AREA	WHAT WENT WRONG	RULE ADOPTED
DETECTION False alarms	Scripted detectors raised false positives (the "24" in an outlet's name; a figure linked to the wrong footnote).	Run cheap scripted scans first, then verify each hit before editing anything.
TESTING Reader's view	Links passed on a local server and were dead when the file was opened from disk.	Test the way the reader opens it. Say plainly when a case couldn't be tested.
ACCESS Blocked source	A key outlet blocked automated fetching; a search snippet showed the names but was not the page.	Escalate in order: API, feed, archive, then the operator's own browser. A snippet is not confirmation. Record any loosened evidence rule as an operator decision.
COST Low yield	Rounds of 30 agents (roughly 150–250k tokens each) sometimes found almost nothing; open sources rarely name units.	Stop a round early when yield is low, and report the low yield as a finding. Prefer targeted gap-filling to large fan-outs.
PLATFORM Safeguard stop	An automated safeguard stopped two Sonnet subagents; they wrote nothing.	Because agents write only to their own files, nothing was lost. Re-run on another model, and log which model did what.
CONTINUITY Scratch files	Work in a temporary scratch folder could have been cleared.	Hand-offs copy working files into the project and include a resume script. Memory notes hold decisions.
CONTINUITY Stale counts	Counts on the cover went stale after the data changed.	Compute cover figures from the built data, and rerun the upstream scripts in a stated order before the final build.

WHEN IT GOES WRONG QUIETLY

Three of the most serious problems (the silent backup override, the duplicate URLs and the generous independence flags) raised no error at all. They surfaced only because a second pass compared the output with something other than itself: a script, a skeptical agent, or the operator asking "walk me through it".

THREE PROMPT ANTI-PATTERNS

Picks that match two lists

"do 1, 2, 5, 6, and 7" fitted two menus at once. Claude checked which was meant. Keep menu numbers unique within a session.

Unbounded quantities

"Use as many agents as you need" collided with the global cap of eight. Claude kept to the cap and said so. Put the cap in `CLAUDE.md`.

Delegating to the unseen

"Let them ask a fable agent for advice" couldn't be confirmed or relayed, because the orchestrator can't see its agents' advisor calls (§5.10).

Doing it yourself

A field guide in three parts: habits for the operator, what a team lead must provide, and a first-week plan with prompts you can paste.

7.1 For the operator

Do

- Set scope through a Claude-run interview, and write the answers down.
- Ask for menus. Say “do not implement” when you only want options.
- State standards as rules with reasons, and let Claude save them.
- Delegate judgment with a standing tie-breaker (“be conservative, don't reach”).
- Attach screenshots to bug reports.
- Ask Claude to “walk me through” any surprising number.
- Say when to pause, and have Claude write a hand-off.

Don't

- Skip the scoping questions. If the dialog closes, reopen it.
- Assume “as many agents as you need” fits your machine.
- Let agents edit shared files.
- Treat a script's “PASS” as proof that the sources are right.
- Ship unfinished analysis. Park it, and say where.
- Leave rules only in chat. They won't survive a new session.
- Let a low-yield round run on because it's already running.

7.2 For the team lead: what to provide

- ☐ **A machine budget.** Decide the agent cap and who may raise it. This project's cap went 2, 8, 4, 8, 12, on a machine that was also running other heavy work.
- ☐ **A named operator with decision rights** over scope, evidence bar and naming, available for short questions over several days.
- ☐ **Protected-data rules**, written before research begins: who may never be named, what may not be shared, where files may live.
- ☐ **API keys and quotas** for search, social-media and satellite sources, kept in the user's home folder and never in the repository.
- ☐ **A release gate that isn't the operator:** a privacy scan, a fail-closed checker and an independent critique.
- ☐ **Encryption for sensitive outputs.** This project built a tool to lock the file behind a password.

- ☐ An expert reader for the subject. The checks in this paper verify traceability, not truth.

7.3 For the agents you launch

Brief	Constrain	Check
• Open with who asked and why. • Require every source to be opened; never cite from memory. • Forbid invented URLs, names, dates and quotes. • Require a verbatim excerpt for every claim.	• One agent, one output file, one folder. • No browser tools; timeouts on downloads. • Report “not researched”, not “none found”, when the budget runs out. • Cap concurrency across the machine.	• Script checks first, agent verdicts second. • A different, skeptical agent verifies. • High-stakes claims get a red-team. • Scripts, not agents, apply changes.

7.4 Day 0: set up the ground

- ☐ Create a project folder and open Claude Code in it. Keep `data/` (evidence), `report/` (sources for the deliverable) and a build script separate from the output.
- ☐ Write a short project `CLAUDE.md` : what the deliverable is, what is a build product (never edited by hand), and what is deprecated.
- ☐ In your global `~/.claude/CLAUDE.md` , set the **machine-wide agent cap**.
- ☐ Fix the layout for agent outputs: `data/raw/<round>/` → `data/verified/` , with backups in `data/backups/` , away from the loader.
- ☐ Confirm which search, fetch and browser tools are available.
- ☐ Create `DECISIONS.md` : date, decision, reason.

7.5 The first day, then the first week

WHEN	DO	PRODUCES
Hour 0	Paste the opening prompt (T1). Answer the scoping questions.	Scope note; first decisions logged
Hour 1	Claude surveys the sources in every relevant language and proposes 6–12 research tracks.	Landscape page; track list

WHEN	DO	PRODUCES
Hours 2–4	Research rounds 1 and 2: agents write raw files, scripts check excerpts, verifiers verify.	Graded, verified claims
Hours 4–6	Set a <code>/goal</code> for coverage (T2). A “completeness critic” lists gaps; targeted rounds follow.	Gap list; targeted rounds
End of day 1	Draft prose with claim markers; build; run the prose checker.	A footnoted first report
Day 2	Ask what is hard to read (“how can this be improved?”). Expect a redesign.	A navigable product
Days 2–3	Structure: incident records, people, units. Plan first, approve in one line, then run.	Record database
Days 3–4	Split work into sessions with owners; agree the build lock. Get an independent critique; refine once.	Layers; critique fixes
Day 5	Audit counts, glossary and sources; a final update pass; write the hand-off.	Release candidate

7.6 Prompt templates

Fill in the brackets. The originals are in §4 and Appendix B.

T1 • OPENING

- › I want you to create a [deliverable, e.g. an HTML report] from the perspective of [role] examining [subject]. I want you to identify [entities: people, units, incidents...]. This initial [deliverable] should focus purely on [facts]. Cite every fact as [footnotes]. Review the work for factual correctness and eradicate hallucinations. Ask me questions to design your investigation and research. Orchestrate as many agents as you need, within the agent cap in CLAUDE.md.

T2 • GOAL

- › `/goal` make sure to look at all [source types and languages] including [partisan or hard-to-reach sides]. Keep working until you feel you've reasonably exhausted the available [open-source] information.

T3 • MENU

- › How can we [improve / extend] this? What gaps exist that we might address? Do not implement anything, I'm just asking.

T4 • PICK

- › Use your best judgment on the open decisions. Do [numbers]. Ask me questions if a decision would be hard to reverse.

T5 · PLAN FIRST

- › Plan and design a process to [goal]. Do not implement until I give the go ahead. (then, after the plan:) [decision 1], [decision 2], ... [decision N]. Search as much as you need.

T6 · STANDARD

- › The priority is [accuracy]. Be more conservative with the information if things are hard to determine from [the sources]. Do not reach. Highlight discrepancies; don't hide things and don't overstate them. Save this as a rule.

T7 · COORDINATION

- › There are [N] other Claude sessions working in this codebase. Coordinate with them so you don't conflict. [Wait until X finishes before Y.] While you wait, think about [design question].

T8 · REDESIGN

- › I want you to redesign [artifact] from the ground up. It was made iteratively and is not coherent. Keep the facts. Help the reader understand [A], see [B], and explore [C]. Balance narrative with data exploration. Follow best practice like [standard], but don't feel constrained by it. Design first, then get an independent critique of screenshots, then refine once.

T9 · PAUSE

- › Let's pause. Create a handoff for the next session. Make sure all working files are in the project and the work can be resumed with one command.

7.7 From worked example to reusable kit

On day 5 the operator asked Claude to “distill the research and production process of the edition report into a reusable process”. The result is a private repository, `auto-researcher`, containing the 17-stage process, operating rules, brief templates, data schemas, dependency-free Python scripts, a minimal edition builder and a synthetic example that builds offline. It contains no data from this project. Without access, Figure 3 and the rules in §6 are the substance of the kit.

WHAT THE KIT COVERS

Ported and tested on the example: scaffold, excerpt checker, promotion, source normalization, corroboration, apply-edits, record checks, prose apply, link validation, scout, both builders and both checkers. **Documented only:** satellite analysis, person photo cards, biography and glossary pipelines, casualty-figure reconciliation, social-media search and the file-lock tool.

Limits and cautions

What this case study doesn't show, and what a careful reader should keep in mind.

Verification is procedural, not final

The checks prove that citations resolve, quotes appear on their pages and numbers match the cited claims. They don't prove the sources are accurate. The verifiers came from the same model family as the researchers, so errors may be correlated. Human spot-checks, such as the INC-0154 walk-through, were essential but not systematic.

No baseline

Nobody did the same investigation without Claude, so any "x times faster" claim would be speculation. What is documented is what was built, how, and where it failed.

Cost was not measured

Output tokens (25.0 million) are reported, but no dollar figure. Measuring cost means applying current rates to the token counts in the transcripts, including cache reads. Large fan-out rounds were expensive for what they found, and the operator later shifted to targeted work.

"Corrected" is not an error rate

Verifiers changed 635 of 1,820 claim records. A change can be as small as one date or one word, and the labels varied between rounds. The number shows that the verifier was active, not how often the researchers were wrong (Appendix E).

One operator, one subject

The operator had strong judgment about evidence and audience. Teams new to a subject will need an expert reviewer. Several of the best decisions here, such as keeping six grade-A claims and downgrading three, were expert calls.

Sensitive subject, changing platform

The project involved named individuals and alleged atrocities. It was kept private, survivors were never named, and outputs could be encrypted. Separately, safeguards, tools and agent limits changed during the project. Check your organization's data rules, and confirm tool behavior before relying on it.

Session index

44 of the 45 sessions have at least one operator prompt. They are listed in start order; the codes are used throughout.

CODE	STARTED	PROMPTS	OPENING PROMPT (TRUNCATED)
S01	Sep 24 00:12	7	I want you to create a report from the perspective of an human rights investigator exami...
S02	Sep 24 11:02	1	can you list for me the prompts I've used within this project?
S03	Sep 24 11:07	2	this report is so long now it's hard to read. the ToC breaks when I jump to certain sect...
S04	Sep 24 11:29	4	can you add images for named individuals if they exist and can be found? this is interna...
S05	Sep 24 11:33	7	I want you to create a new evidence base database for every named individual. I want you...
S06	Sep 24 11:38	18	What gaps exist in the report and how can we improve it?
S07	Sep 24 14:56	6	the report should use names that reflect the latin spelling found in reports. like in th...
S08	Sep 24 14:57	6	how can we build upon this report? file://~/sudan-research/report/kordofan_rsf_package.h...
S09	Sep 24 17:34	6	resume phase 3"
S10	Sep 24 18:13	1	I'm not happy with the how the analysis tab is designed. I find it hard to understand. P...
S11	Sep 24 19:14	9	I'd like for you to plan and design a process to enrich the current incidents we have wi...
S12	Sep 24 19:27	1	I want you to look at the unfinished sentinel 1 and 2 analysis work we were doing and re...
S13	Sep 24 22:51	1	I want to separate out the glossary from the named individuals
S14	Sep 24 23:07	2	I want you to bring the analysis and any other relevant sections up to date based on the...
S15	Sep 24 23:30	3	I want you to redesign file://~/sudan-research/report/kordofan_rsf_package.html from the...
S17	Sep 25 00:13	4	I want you to polish and refine the design, styling, UI for file://~/sudan-research/repo...
S18	Sep 25 07:54	2	@~/sudan-research/report/kordofan_edition.html" you mentioned in a previous session tha...
S19	Sep 25 08:48	3	I'd like for you to add a dedicated satellite imagery analysis section in the Explore th...
S20	Sep 25 09:02	4	I want to include in the number counts here how many people were reported killed during ...
S21	Sep 25 09:07	4	are there any speech to text tools that address kachaqil?
S22	Sep 25 09:58	6	I want you to go through and update the document throughout based on all of the latest r...
S23	Sep 25 10:04	1	would it be helpful if you could use apify to search on social media to supplement infor...
S24	Sep 25 10:04	7	would it be helpful if you could use apify to search on social media to supplement infor...
S25	Sep 25 10:14	2	I'd like to create a tool that would password protect and encrypt the .html file output,...
S26	Sep 25 10:21	1	I'd like for you to visualize the order of battle and then make it a distinct section in...
S27	Sep 25 14:09	1	when when i select in the chart it doesn't seem to actually all appear in the map
S28	Sep 25 17:34	16	how can we either further visualize and analyze the existing data or expand our research...

CODE	STARTED	PROMPTS	OPENING PROMPT (TRUNCATED)
S29	Sep 25 21:59	3	I want you to incorporate FIRMS into this analysis file://~/sudan-research/report/kordof...
S30	Sep 25 22:05	2	I want you to redesign the cards in the north, south and west kordofan sections to be or...
S31	Sep 25 22:35	1	I want you to ensure that the styling and design is consistent throughout this tool file...
S32	Sep 25 22:49	3	I want you to build out profiles for units. conduct as much research as you need. in add...
S33	Sep 25 22:54	1	I want you to find a better way to represent the information contained in the Responsibi...
S34	Sep 25 23:16	1	I want you to add a background section about why the fighting is happening in Kordofan -...
S35	Sep 25 23:29	1	I want you to examine all charts labels and relevant sections to ensure that the labels ...
S36	Sep 25 23:45	1	I want you to make sure these count are up to date and also to make sure sources, glossa...
S37	Sep 27 20:40	3	how can this be improved?
S38	Sep 27 20:58	8	can you look for SAF, RSF and OSINT channels on Telegram and expand our Telegram explora...
S39	Sep 27 21:02	2	I want you to redesign this page file://~/sudan-research/report/kordofan_edition.html#/u...
S40	Sep 28 11:31	3	please make it so that the source documents count just count all sources used
S41	Sep 28 11:38	1	how much work would it be to add a section looking at SAF activities? I want to keep it ...
S42	Sep 28 18:12	3	how can we improve the edition version?
S43	Sep 28 23:27	5	can you close any of the research gaps? if, what?
S44	Sep 29 09:49	1	I'd like for you to distill the research and production process of the edition report in...
S45	Sep 29 10:53	1	I'd like for you to prepare a case study describing how Claude Code was used to create t...

Every prompt, verbatim

All 165 prompts the operator typed, in time order. A new header marks each switch between sessions.

`/goal` marks the condition text of a goal. `resume prompt` marks two prompts that appear to be session-resume notes. `+ screenshot` means an image was attached within four minutes. System messages, injected skill text, cross-session notices and subagent replies are excluded because the operator didn't write them.

S01 Sep 24

00:12 I want you to create a report from the perspective of an human rights investigator examining the RSF's fighting in the Kordofan region. I want you to identify units, commanders, specific people, specific incidents of human rights abuses, mass killings, torture, rape. This initial report should be focused purely on the facts. I want you to create a html report, citing all facts as footnotes. I want you to review the work to check for factual correctness and to eradicate all hallucinations. Ask me questions to design your investigation and research. Orchestrate as many agents as you need to do the work thoroughly, quickly and effectively

00:13 ask me the questions again please

00:20 `/goal` make sure to look at all sources including arabic language sources and sources from the SAF and RSF sides. keep working until you feel like you've reasonably exhausted the available open source information

07:23 Give me a status update

07:54 I want you to add an always present ToC to the .html report to help the user navigate the report. I also want you to add an elegant, professional and beautiful design to the .html file. Use charts and tables where relevant to give users a high level view of the information.

08:12 update the summary to include round 3 findings. I also want you to add a section specifically about any units or commanders tied to atrocities

08:28 add perpetrators to at a glance

S02 Sep 24

11:02 can you list for me the prompts I've used within this project?

S03 Sep 24

11:07 this report is so long now it's hard to read. the ToC breaks when I jump to certain sections because there is too much content. Please redesign this report for legibility and clarity. If sections need to be a database or a radical redesign, do so

11:20 I want the table of contents to be more obvious and always on if possible

S04 Sep 24

11:29 can you add images for named individuals if they exist and can be found? this is internal use so IP won't matter, this will never be published

11:30 I'd also like for you to include arabic spellings for all commander names. For each person profiled, there should be a table at the top giving basic biographic info and things useful for an osint investigator to know like name variations in different languages, aliases, etc

S05 Sep 24

11:33 I want you to create a new evidence base database for every named individual. I want you to build on the existing information in the report and add a full profile giving a personal, professional and military biography. create a timeline for each named individual on this profile page. cite sources do a verification pass to ensure eradicate all hallucinations

S06 Sep 24

11:38 What gaps exist in the report and how can we improve it?

11:40 there are two other claude sessions working on this project right now, please coordinate with them

11:42 yes, go ahead with both

11:46 yes, downgrade those 3

S05 Sep 24

11:46 go ahead once the other session frees a slot

S06 Sep 24

11:46 yes, look for the Yale report

S05 Sep 24

11:48 ping the images session for an ETA

S06 Sep 24

11:49 keep the other 6 at grade A

S05 Sep 24

11:50 let's save [name], [name] for last, do the less obvious commanders first

S06 Sep 24

11:50 go ahead with the incident map

11:51 yes, download both

12:06 start on the thin months for Jan-Apr 2025

12:18 add them to the section prose

S05 Sep 24

12:18 keep going through the batches

S06 Sep 24

13:39 do the sweep for June, August and September 2025

13:45 add them to the section prose

14:00 start the West Kordofan corroboration pass

S05 Sep 24

14:03 keep going

S06 Sep 24

14:05 we can use more than 2 agents at a time. let's set the limit to 8 total

S04 Sep 24

14:06 find photos for the people still missing them

14:07 also find Arabic sources for the four missing ones

S06 Sep 24

14:07 apply the results when they come in

14:08 then rebuild and summarize when all batches finish

14:40 reword the [name] claim and split out the Ghubaysh arrests

S05 Sep 24

14:40 link the profiles page from the main report

S06 Sep 24

14:50  when i click on a name on the left it doesn't take me to the profile.

S07 Sep 24

14:56  the report should use names that reflect the latin spelling found in reports. like in this example, it should not transliterate in a way that is not represented in the cited sources

S08 Sep 24

14:57 how can we build upon this report? file://~/sudan-research/report/kordofan_rsf_package.html

14:58 I want to do 2 and 3

S07 Sep 24

15:21 resume prompt Check name-spelling agents' outputs in scratchpad/names/out and continue the Latin-spelling fix.

15:22 can you find a way around sudan tribune's blocked automated fetching? i think they're a key source

15:25 I'm clicking it, it won't verify

S08 Sep 24

15:25 you can use more searches

S07 Sep 24

15:26 use my chrome

15:27 ok extension connected, try again

S08 Sep 24

16:20 go ahead and finish the rebuild

16:20 actually let's pause, I'm going to relocate

16:21 create a handoff for the next session

S09 Sep 24

17:34 resume phase 3"

17:52 run the companions through command responsibility review

17:55 approve the three wording checks

17:57 yes, raise the Dilling claim to A

18:06 raise the 3 affected by today's regrades

18:09 add the attribution notes to the analysis prose

S10 Sep 24

18:13 I'm not happy with the how the analysis tab is designed. I find it hard to understand. Please rework it so that the analysis is clear and able to be immediately understood. Keep the nuance.

S11 Sep 24

19:14 I'd like for you to plan and design a process to enrich the current incidents we have with further research to build out incident profiles and a database of incidents in Kordofan that could be a human rights abuse or atrocity. Do not implement the work until I give the go ahead

19:16 continue working

19:19 Descriptive only, RSF+allies, parent campaigns, new tab, feeding back, name killed victims. unleash the search budget, search as much as you need.

S12 Sep 24

19:27 I want you to look at the unfinished sentinel 1 and 2 analysis work we were doing and remove it from the final product. save it for us to finish later

S11 Sep 24

19:27 I actually want to cap this at 4 agents

19:32 keep working through the rest of the phases once the review is done

19:32 you're allowed to use 8 agents again

20:01 keep going, no need to check in

21:00 update me on what's been done so far, what you've found and next steps

21:02 leave out weak new incidents, keep it in the database. note that there are weak incidents

S13 Sep 24

22:51 I want to separate out the glossary from the named individuals

S14 Sep 24

23:07 I want you to bring the analysis and any other relevant sections up to date based on the latest research

S15 Sep 24

23:30 I want you to redesign file:///~/sudan-research/report/kordofan_rsf_package.html from the ground up. it was made in an interactive process and is not coherent or unified. I want you to recreate it, keeping the facts and information, in a way that helps the reader understand the situation, who is responsible, and be able to easily explore the underlying data. It should balance narrative with data exploration. It should follow best practices like those outlined by Tufte, but do not feel constrained by it. Strive for beauty, elegance, legibility, understandability. Help the user understand the information - show linkages, patterns and highlight connections. Make this a delight for the user to read and experience but keep it informative - something the user can come back to to use as a tool. Use your full creativity

S14 Sep 24

23:32 apply the card edits, that rule is fine, have it as person unity or party

S15 Sep 24

23:34 do it after your review

S17 Sep 25

00:13 /goal I want you to polish and refine the design, styling, UI for file:///~/sudan-research/report/kordofan_edition.html

I like it, I want to polish and improve it. Continue working until you feel like it is production ready and can be published to a wide audience. Pay attention to close details, especially within tables, making sure that small details like headers for tables freeze as people scroll and other quality of life design elements are all in place. ask me if you have any questions. there's another session working in this codebase, coordinate with it

00:17 keep working without my input

00:49 Check the dates yourself and correct as needed. Keep working

00:55 Examine the report for any UI gaps

S18 Sep 25

07:54 @"/sudan-research/report/kordofan_edition.html"
you mentioned in a previous session that there are broken links - how many broken links exist in this

08:19 fix the two

S19 Sep 25

08:48 I'd like for you to add a dedicated satellite imagery analysis section in the Explore the Evidence section and build it out. The goal should be to use sentinel 1, 2 and wayback imagery to illuminate the incidents. Use your best judgment to determine which incidents can be explored using the available data. I want you to push the boundaries a bit here, but keep it legible, the analysis understandable. use GEE and other tools as needed.

S15 Sep 25

08:55 let the other session do the build, we're done here

S20 Sep 25

09:02 + screenshot I want to include in the number counts here how many people were reported killed during the period of investigation

09:06 run the reconciliation pass

S21 Sep 25

09:07 are there any speech to text tools that address kachaqil?

09:09 can you check the two meta options?

09:11 can you scope out implementing Meta Omnilingual ASR but do not test it yet on audio clips, I have other things running at the moment

09:11 can we use it to detect Kaqchikel clips and segment and separate it from the spanish?

S20 Sep 25

09:22 first walk me through INC-0154. then address INC-0129 and the Misfiled claims.

09:25 i'll take your recommendations here, use your best judgment and implement fixes for all three

S19 Sep 25

09:46 what do you mean 5 months from done?

09:47 fix the wrong locatoin in the database. acquire the missing imagery if it exists

S22 Sep 25

09:58 I want you to go through and update the document throughout based on all of the latest research and findings. use as many agents as you need file://~/sudan-research/report/kordofan_edition.html

S23 Sep 25

10:04 would it be helpful if you could use apify to search on social media to supplement information?

S24 Sep 25

10:04 would it be helpful if you could use apify to search on social media to supplement information with an OSINT approach?

10:06 I've developed a set of tools here, evaluate if you can use it for this work.
~/Documents/social_monitor

S22 Sep 25

10:06 I want you to be proactive and work on other things if you are able to

S24 Sep 25

10:09 end run_all.sh. that's fine to use the remaining budget. let's do pilot first. create a script to account for telegram. address the searching by incident issues as well

S22 Sep 25

10:10 I'll allow up to 12 now

S25 Sep 25

10:14 I'd like to create a tool that would password protect and encrypt the .html file output, similar to the one found here [path to another project]

I don't want the right click protections, people should be able to select text and copy as they choose.

I do want to password protect it and encrypt the .html so that no one can read the information even if they find it in a github repo

S22 Sep 25

10:16 I want you to use your best judgment to address those notes. The priority is accuracy - be more conservative with the information if things are hard to determine from the open source - do not reach

S26 Sep 25

10:21 I'd like for you to visualize the order of battle and then make it a distinct section in the Explore the Evidence rail. The design should effectively communicate the order of battle, command structures if we know them, and be able help the user visualizer and understand the units and how they relate to each other or don't relate to each other file://~/sudan-research/report/kordofan_edition.html#/read/order-of-battle

S25 Sep 25

10:22 can you remove from git history?

S24 Sep 25

10:24 keep me updated. be proactive if you can work on something

10:25 make sure the work can be paused and resumed

S22 Sep 25

10:27 I want you to update the glossary too based on all the new information

10:34 @"/sudan-research/report/kordofan_rsf_report_linear.html" @"/sudan-research/report/kordofan_rsf_report.html"
just fyi these two files are deprecated, please save that as a memory and update claude.md to reflect that

S24 Sep 25

10:59 add the verified claims, do not move it into scope, save the 87 cited post urls to wayback

implement next steps. I also want you to evaluate if it is worth scaling up or addressing more incidents with it

11:23 skip wayback for now

S27 Sep 25

14:09  when when i select in the chart it doesn't seem to actually all appear in the map

S28 Sep 25

17:34 how can we either further visualize and analyze the existing data or expand our research? What investigative gaps exist that we might be able to address? do not implement anything I am just asking

19:31 use your best judgment to address the open decisions. I want you to implement unit attribution and the easy work first. do 1, 2, 5, 6, and 7

20:22 update the tdate and the scope sentence

20:30 do the command link review, restart telegram

21:06 where are we at with telegram

21:07 yeah, triage the new candidates

21:08 does it make sense to update the report text again?

21:08 let's move forward with your recommendation

21:15 I want you to review the all of the graphics and charts and determine if they are worth an update based on the new information

21:19  i want to be able to toggle on and off the control lines

21:27 any other next steps to address?

21:31 please do 1, 2, 3 5 and 6

21:34 so text and everything else is all up to date?

21:34 update sources to reflect x and telegram sources

21:36 add filtering sorting and group the sources for me. make it useful for someone who might be looking for sources to do further research

21:39  make the filters more compact. spend some time to figure out the best UI for this. wherever else this type of filtering ui shows up, I also want you to make it more compact

S29 Sep 25

21:59 I want you to incorporate FIRMS into this analysis file://~/sudan-research/report/kordofan_edition.html#/satellite

S30 Sep 25

22:05 I want you to redesign the cards in the north, south and west kordofan sections to be organized in the existing timeline format or with a toggle option to be organized into groups based on categories of types of attacks. there are two other sessions working in those codebase. one is working on the cards right now, coordinate so that you don't conflict. wait until it's done with updating the cards to implement any changes. while you wait, think of how you can make this design elegant beautiful and useful

S29 Sep 25

22:14 I want you to stack the FIRMS detections with wayback imagery for context

S30 Sep 25

22:28  the numbers in the sections for patterns of abuse don't make sense to me

S31 Sep 25

22:35  I want you to ensure that the styling and design is consistent throughout this tool file://~/sudan-research/report/kordofan_edition.html

run a deep QA on the tool to prepare it for use by investigators

S32 Sep 25

22:49 I want you to build out profiles for units. conduct as much research as you need. in addition to this, make sure that all commanders have a profile and all the information is up to date based on the latest ressearch.

S33 Sep 25

22:54 I want you to find a better way to represent the information contained in the Responsibility section. style each named commander and allow for grouping and sorting in a similar way that the "The Three Fronts" section organizes its information. Take inspiration and localize it to make sense for the Responsibility section

S29 Sep 25

22:59 the crops are too small to see. I also want to be able to toggle off the notation

S34 Sep 25

23:16 I want you to add a background section about why the fighting is happening in Kordofan - what is the strategic and economic value? make sure it is well researched and designed

S35 Sep 25

23:29 /goal + screenshot I want you to examine all charts labels and relevant sections to ensure that the labels are clear, simply written and legible. For example, the labels and titles in this screenshot do not read well and feels overly complicated to read. You are only adjusting the text, not the design or graphics themselves

S32 Sep 25

23:42 resume prompt Check research agent progress (data/bios/upd/raw, data/units/raw) and continue the unit/commander profile pipeline: quote-check, verify, apply, merge, rebuild.

23:43 write them into the chapters where they are relevant

S36 Sep 25

23:45 + screenshot I want you to make sure these count are up to date and also to make sure sources, glossary and claims sections are up to date reflecting the latest information. coordinate with the other sessions

S37 Sep 27

20:40 how can this be improved?

20:44 go ahead with 1-8. use your best judgment in implementing this. ask me questions to help guide the design and fixes

20:45 continue

S38 Sep 27

20:58 can you look for SAF, RSF and OSINT channels on Telegram and expand our Telegram exploration.

S39 Sep 27

21:02 I want you to redesign this page file://~/sudan-research/report/kordofan_edition.html#/units

21:03 file://~/sudan-research/report/kordofan_edition.html#/read/order-of-battle

S38 Sep 27

21:06 how far along are we with the crawl?

21:11 resume

21:13 add the new channels to the list and run that in an additoinal phase

21:17 can you give me a telegram channel that forwards a lot?

21:24 keep me updated on the telegram crawling

22:02 can we do a preliminary analysis

22:16 update me on what's been found since the last update

S40 Sep 28

11:31  please make it so that the source documents count just count all sources used

11:32 for the total counts, does that mean we have to do the analysis for 20 incidents to update the number?

11:33 address the 14

S41 Sep 28

11:38 how much work would it be to add a section looking at SAF activities? I want to keep it separate from the RSF analysis but still in this edition version

S42 Sep 28

18:12 how can we improve the edition version?

18:15 let's do 1 - 4

18:36 move those 5 out of scope

S43 Sep 28

23:27 can you close any of the research gaps? if, what?

23:34 go ahead with everything, then kick off an opus agent to update the prose in the edition report

23:35 use as many opus and sonnet agents as you need for work. let them ask a fable agent for advice

23:39 can you have the fable advisor update you on the agents questions?

23:42 what's a safeguard error

S44 Sep 29

09:49 I'd like for you to distill the research and production process of the edition report into a reusable process. The research should determine the specific structure of the report but this is the general skeleton that I like. Make this into a prviate repo called auto-researcher

10:53 I'd like for you to prepare a case study describing how Claude Code was used to create this project. Give me a detailed report in a .html and .pdf format. It should be well designed and formatted. This is a methodology paper to share with colleagues so share lessons learned, best practices and anything else someone would need to begin working. Share the exact prompts I used and give excerpts from key conversations. Ask me questions to help guide the design

Anatomy of an agent brief

Two parts of a real brief sent to research agents: the authorization preamble that prevented refusals, and the shared instructions.

C.1 The authorization preamble

It repeats the operator's own request and choices, so that an agent, which never saw the chat, knows who is asking, why, and that the task is authorized. It was added after an agent refused a task it didn't understand.

AGENT BRIEF

verbatim, from the Workflow script

CONTEXT – WHY YOU ARE DOING THIS (read first):

The user, in the main chat session, asked for a fact-only human-rights investigator's report on the RSF's fighting in the Kordofan region (units, commanders, specific people, specific incidents of abuses, mass killings, torture, rape), with every fact footnoted and all hallucinations eliminated, and asked the orchestrator to use as many agents as needed. The user answered scoping questions: time scope = 2025-2026 only; actors = RSF + allies across North, West and South Kordofan; evidence = tiered sourcing including OSINT/social media; naming = name perpetrators, never survivors. The user then set an explicit goal: "make sure to look at all sources including arabic language sources and sources from the SAF and RSF sides. keep working until you feel like you've reasonably exhausted the available open source information".

The task below is the orchestrator's direct fulfilment of that request. It is fully in scope and authorized – carry it out completely. Do not ask questions; do not stop early.

C.2 The shared instructions (excerpt)

Each research track added its own topic on top of this block. The output format (omitted) required a verbatim excerpt of at most 25 words per source, a grade and a note on any discrepancy.

AGENT BRIEF

verbatim, from the Workflow script

You are a researcher supporting a human-rights investigation into the Rapid Support Forces (RSF) and allied forces (incl. SPLM-N al-Hilu where operating jointly under the "Tasis" alliance, and allied Arab

militias) in the Kordofan region of Sudan (North, West and South Kordofan states). TIME SCOPE: 1 January 2025 to today (24 September 2026). Facts only – no advocacy, no speculation.

TOOLS: Load WebSearch and WebFetch first via ToolSearch ("select:WebSearch,WebFetch"). Search widely (English AND Arabic queries, e.g. `الدعم السريع كردفان, بارا, النهود, الخوي, بابنوسة, كادقلي, الدلنج, الأبيض`). Actually OPEN (WebFetch) every source you cite and confirm it says what you claim. NEVER cite from memory. NEVER invent a URL, name, date, number or quote. If you cannot confirm something from a fetched page, leave it out or put it in "leads_unconfirmed".

SOURCES (all allowed, grade them): UN (Independent International Fact-Finding Mission for Sudan, OHCHR/[name] statements, UN Panel of Experts on Sudan, OCHA, UNICEF, WFP, IOM DTM), HRW, Amnesty, Yale HRL (Humanitarian Research Lab), ACLED, Sudan War Monitor, Centre for Information Resilience (Sudan Witness), Bellingcat, Sudan Tribune, Radio Dabanga, Reuters, AP, AFP, Al Jazeera, BBC, Guardian, NYT, CNN, Emergency Lawyers, Sudan Doctors Network, local resistance committees, government statements (flag as partisan), RSF statements (flag as partisan), OSINT/social media incl. RSF-posted videos (state verification status: geolocated by whom, or unverified). Sanctions designations: US Treasury OFAC, UK FCDO/OFSI, EU Council, UN 1591 Committee.

NAMING RULE: Name perpetrators/commanders where a credible source attributes conduct to them. NEVER name survivors of rape/sexual violence or torture. Name killed victims only if the source already published the name (prefer not to).

CONFIDENCE GRADES: "A" = 2+ independent credible sources or UN/major-NGO documented investigation; "B" = one credible source (major outlet/NGO/UN statement) or credible local source corroborated by OSINT; "C" = single local/partisan/unverified social media source.

Claude Code features used

Which capabilities carried which part of the method, with usage counts from the transcripts.

FEATURE	USED FOR	SEE
<code>AskUserQuestion</code>	Structured multiple-choice interviews (13 uses in the main sessions)	\$3.1, \$5.1
<code>Subagents (Agent)</code>	Research, verification, critique, drafting (456 launches)	\$3.2–3.3
<code>Workflow</code> <code>scripts</code>	Phased fan-out: research, then verify (17 runs)	\$3.2, App. C
<code>/goal</code>	Keep a session working to a stated finish line (4 goals)	\$3.6
<code>SendMessage</code> , <code>idle notices</code> , <code>ListAgents</code>	Coordination between sessions (199 messages)	\$3.5, \$5.2

FEATURE	USED FOR	SEE
<code>CLAUDE.md</code> and memory notes	Standards and decisions across sessions (about 30 notes)	\$3.7
Built-in browser pane	Previewing the report, checking layouts and console errors (over 1,000 calls)	\$5.11
Claude in Chrome	Reading a source that blocked automated fetching; read-only, at the operator's request	\$5.5
Skills	Workflow-authoring and page-design guidance, loaded on demand	—
Advisor, model mix	A second model consulted on hard calls; Opus and Sonnet subagents	\$5.10
<code>Monitor</code> , <code>TaskStop</code>	Watching long background jobs; stopping them cleanly	\$5.6

Tool calls across the project

Main sessions plus 703 subagent transcripts. `Bash` dominates because Claude wrote and ran scripts for nearly everything; direct file edits were an order of magnitude rarer.

<code>Bash</code> scripts, builds, checks		14,085
<code>WebFetch</code> open a source page		6,837
<code>WebSearch</code>		5,813
<code>Read</code>		1,280
<code>Write</code>		830
<code>Edit</code>		672
<code>Agent</code> subagent launches		456
<code>ToolSearch</code>		409
<code>SendMessage</code> between sessions		199

Method and definitions

Data freeze. Transcript data was frozen at Sep 29, 11:00 local (18:00 UTC). `prep_data.py` ignores anything later, so this paper's own session doesn't change the counts on a re-run.

Sources. Claude Code transcripts (`.jsonl`) for this project folder: 45 main sessions and 703 subagent transcripts. Also the project's memory notes, the edition checker's output, and the `auto-researcher` kit documents.

Operator prompts. User messages with human-typed text. Excluded: injected skill text, compaction summaries, cross-session notices, subagent replies, “request interrupted” markers and an automated “resume” instruction. Screenshots are folded into the preceding prompt.

Prompt length. Whitespace-separated words in the prompt as typed. Median 10; 78% under 25 words.

Time. Local time, UTC−7 (matches the operator's memory notes).

Agent launches. `Agent` and `Workflow` calls issued by main sessions (473 = 456 + 17). Subagents' own tool calls are included in the tool chart but not in the daily chart.

Output tokens. Sum of `output_tokens` on assistant messages across main and subagent transcripts. Cache reads (billions) are deliberately not reported as work.

Concurrency. The number of distinct main sessions with any message in the same 30-minute window. The maximum was six.

Claim verdicts. Of 1,820 claim records in the verified store (excluding backups): verified 805, corrected 635, confirmed 373, rejected 7. Labels differ between rounds, so “corrected” is read here as “the verifier changed the claim”, not as an error rate.

Product counts. From the edition checker: 1,813 claims, 277 incidents, 85 people, 437 glossary entries, 13 sections, 19.01 MB, PASS. Source and publisher counts are from the Sep 26 audit and may have changed.

Code size. About 16,400 lines of Python, JavaScript and CSS (65 Python scripts plus the edition's JavaScript and CSS), excluding data and backups.

Quotations. Checked by the build script against the transcripts. Names and name-revealing details are redacted; cuts are marked [...]; emoticons are removed from Claude's replies.

Typefaces. Inter, Source Serif 4 and JetBrains Mono, all under the SIL Open Font License, embedded so the file works offline.

Prepared with Claude Code from the project's own records. Working paper, September 2026. To reproduce this document, run `python3 prep_data.py && python3 build_case_study.py && python3 make_pdf.py` in `case_study/`.